

18-661 Introduction to Machine Learning

Reinforcement Learning

Spring 2025


ECE – Carnegie Mellon University

1. MDP Recap
2. Bellman Equations
3. Finding the Optimal Policy
4. Q-learning
5. Course Summary

MDP Recap

Overview of Reinforcement Learning

In RL, an agent learns through interaction with the (unknown) environment.



Objectives of RL

We choose the actions that **maximize the expected total reward**:

$$R(T) = \sum_{t=0}^T \mathbb{E} [r(a(t), s(t))], \quad R(\infty) = \sum_{t=0}^{\infty} \mathbb{E} [\gamma^t r(a(t), s(t))].$$

Infinite-horizon discounted reward. We *discount* the reward at future times by $\gamma < 1$ to ensure convergence when $T = \infty$. The expectation is taken over the probabilistic evolution of the state, and possibly the probabilistic reward function.

A **policy** tells us which action to take, given the current state.


- **Deterministic policy:** $\pi : S \rightarrow A$ maps each state s to an action a .
- **Stochastic policy:** $\pi(a|s)$ specifies a probability of taking each action $a \in A$ given state s . We draw an action from this probability distribution whenever we encounter state s .

Transition Matrices

A Markov chain can be represented with a **transition matrix**:

$$P = \begin{bmatrix} p_{1,1} & p_{1,2} & \dots & p_{1,n} \\ p_{2,1} & p_{2,2} & \dots & p_{2,n} \\ \vdots & \vdots & \ddots & \vdots \\ p_{n,1} & p_{n,2} & \dots & p_{n,n} \end{bmatrix}$$

Each entry (i, j) is the probability of transitioning from state s_i to state s_j .



$$P = \begin{bmatrix} 0 & 1 & 0 & 0 \\ 0.5 & 0 & 0.2 & 0.3 \\ 0.9 & 0 & 0 & 0.1 \\ 0 & 0 & 0.8 & 0.2 \end{bmatrix}$$


Markov Decision Process

Markov decision processes make state transitions dependent on user actions


Markov processes represent transitions from state $s(t)$ to $s(t + 1)$, but they do not include actions taken by the user. We now need to introduce a set of transition matrices:

$$p_{i,j} \rightarrow p_{i,j}^a = P(s(t + 1) = s_j | s(t) = s_i, a(t) = a),$$

where a is a given action.



Action 1



Action 2

Evaluating a Policy: The State-Value Function

The **state-value function** of a given policy $\pi : S \rightarrow A$ gives its expected future reward when starting at state s :

$$V_{\pi}(s) = \mathbb{E}_{a \sim \pi(s), P, r} \left[\sum_{t=0}^{\infty} \gamma^t r(a(t), s(t)) \mid s(0) = s \right].$$

The expectation may be taken over a stochastic policy and reward as well as the Markov decision process (MDP) of the state transitions.

- The action $a(t)$ at any time t is determined by the policy, $\pi(s(t))$.
- Due to the Markov property of the underlying MDP, the optimal policy at any time is only a function of the last observed state.

The Action-Value Function

Consider a variant of $V_\pi(s)$, the Q function

$$Q_\pi(a, s) = \mathbb{E} \left[r(a, s) + \underbrace{\sum_{t=1}^{\infty} \gamma^t r(\pi(s(t)), s(t))}_{\text{use policy } \pi} \mid s(0) = s, a(0) = a \right]$$

The Q function gives the expected value of taking an action a at time 0, and then following a given policy π .

Bellman Equations

Evaluating the State-Value Function

Let the sequence of received rewards be r_0, r_1, \dots following policy π :

$$\begin{aligned} V(s) &= \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t r_t \mid s_0 = s \right] \\ &= \mathbb{E} [r_0 + \gamma r_1 + \gamma^2 r_2 + \gamma^3 r_3 + \dots \mid s_0 = s] \\ &= \mathbb{E} [r_0 \mid s_0 = s] + \gamma \mathbb{E} [r_1 + \gamma r_2 + \gamma^2 r_3 + \dots \mid s_0 = s] \\ &= \mathbb{E} [r(s)] + \gamma \mathbb{E}_{s' \sim P(s_1 | s_0 = s)} \mathbb{E} [r_1 + \gamma r_2 + \gamma^2 r_3 + \dots \mid s_0 = s, s_1 = s'] \\ &= \mathbb{E} [r(s)] + \gamma \mathbb{E}_{s' \sim P(s_1 | s_0 = s)} \mathbb{E} [r_1 + \gamma r_2 + \gamma^2 r_3 + \dots \mid s_1 = s'] \\ &= \mathbb{E} [r(s)] + \gamma \mathbb{E}_{s' \sim P(s_1 | s_0 = s)} V(s') \end{aligned}$$

The value function can be decomposed into two parts

- immediate reward $\mathbb{E} [r(s)]$
- discounted value of at the successor state $\gamma \mathbb{E}_{s' \sim P(s_1 | s_0 = s)} V(s')$

this defines a recursive relationship between the state values!

Bellman Equation

$$\forall s : V(s) = \mathbb{E}[r(s)] + \gamma \mathbb{E}_{s' \sim P(s_1|s_0=s)} V(s')$$

Put this in a matrix form using the probability transition matrix \mathbf{P} induced by the policy, and

$$\mathbf{V} = \begin{bmatrix} V(s_1) \\ V(s_2) \\ \vdots \\ V(s_n) \end{bmatrix}, \quad \mathbf{r} = \begin{bmatrix} \mathbb{E}[r(s_1)] \\ \mathbb{E}[r(s_2)] \\ \vdots \\ \mathbb{E}[r(s_n)] \end{bmatrix}$$

we have the Bellman equation:

$$\mathbf{V} = \mathbf{r} + \gamma \mathbf{P} \mathbf{V}$$

Bellman Equation

$$\mathbf{V} = \mathbf{r} + \gamma \mathbf{P}\mathbf{V}$$


- We can solve the Bellman equation directly by matrix inversion:

$$(\mathbf{I} - \gamma \mathbf{P})\mathbf{V} = \mathbf{r}$$

$$\mathbf{V} = (\mathbf{I} - \gamma \mathbf{P})^{-1} \mathbf{r}$$

- Computational complexity is $O(n^3)$ for n states
- Try iterative methods for large MRPs, e.g. dynamic programming, monte-Carlo evaluation, temporal-difference learning.

Example



$$P = \begin{bmatrix} 0 & 1 & 0 & 0 \\ 0.5 & 0 & 0.2 & 0.3 \\ 0.9 & 0 & 0 & 0.1 \\ 0 & 0 & 0.8 & 0.2 \end{bmatrix}$$

Assume $r(s_1) = 1$, $r(s_2) = 1$, $r(s_3) = 1$, $r(s_4) = 10$, and $\gamma = 0.9$, then the value function is

$$V = (I - \gamma P)^{-1} r = \begin{bmatrix} 21.5896 \\ 22.8773 \\ 21.2656 \\ 30.8674 \end{bmatrix}$$

Verify: $V(s_3) = r(s_3) + 0.9[0.9V(s_1) + 0.1V(s_4)]$



Bellman Equations for MDP

Bellman equation for state-value function:

$$V_{\pi}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) Q_{\pi}(a, s),$$

Bellman equation for action-value function:

$$Q_{\pi}(a, s) = \mathbb{E}[r(a, s)] + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V_{\pi}(s'),$$




Bellman equation for state-value function:

$$V_{\pi}(s) = \mathbb{E}[r(a, s)] + \gamma \sum_{a \in \mathcal{A}} \pi(a|s) \sum_{s' \in \mathcal{S}} P(s'|s, a) V_{\pi}(s')$$

Example

Consider the same maze example as earlier, where $\gamma = 1$ and

- we receive a reward of $+1$ if we reach $[4,3]$ and -1 if we reach $[4,2]$, which terminate the epoch; -0.04 for taking an action;
- The policy is plotted in arrows; where if there are two arrows in a cell, both directions take equal probability.
- One can certainly use the Bellman equation to solve the value function, but for this example, we can also find it by “dynamic programming” from the end state.



Finding the Optimal Policy

Optimizing the Action-Value Function

- Suppose we know the “optimal” Q-function, $Q^*(a, s)$, for each action a and state s .
- Then, we maximize the state-value function by choosing the action $a^*(s)$ at state s so as to maximize $Q^*(a, s)$.

$$V^*(s) = \max_a Q^*(a, s)$$

But this gives us π^* !

$$\pi^*(s) = \arg \max_a Q^*(a, s)$$

$$V^*(s) = V_{\pi^*}(s)$$

greedification of Q^* yields the optimal policy!

- For tabular MDPs, there is always a *deterministic* optimal policy (not necessarily unique).

Bellman's Optimality Equation

The optimal value functions are recursively related by the Bellman optimality equations:

$$V^*(s) = \max_a Q^*(a, s)$$

$$Q^*(a, s) = \mathbb{E}[r(a, s)] + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V^*(s')$$

Putting them together:

$$V^*(s) = \max_a \left[\mathbb{E}[r(a, s)] + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V^*(s') \right]$$

$$Q^*(a, s) = \mathbb{E}[r(a, s)] + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) \max_a Q^*(a, s')$$

These equations define recursive relationships that allow us to find the optimal value functions.

Solving the Bellman Optimality Equation

Bellman operator: one-step look-ahead

$$\mathcal{T}(Q)(a, s) = \mathbb{E}[r(a, s)] + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) \max_a Q(a, s')$$

then the Bellman's optimality equation says Q^* is the (unique) fixed point of \mathcal{T} :

$$Q^* = \mathcal{T}(Q^*)$$

- Bellman Optimality Equation is non-linear
- There is in general no closed form solution
- Many iterative solution methods
 - Value Iteration
 - Policy Iteration
 - Q-Learning
 - more...

Value Iteration

Initialize $V(s)$ for each state s . Suppose we know the reward function $r(a, s)$ and transition probabilities $P(s'|s, a)$.

- Update $Q(a, s)$:

$$Q(a, s) \leftarrow \mathbb{E}[r(a, s)] + \gamma \underbrace{\sum_{s' \in \mathcal{S}} P(s'|s, a) V(s')}_{\text{Estimate of expected future reward}}$$

- Update $V(s) \leftarrow \max_a Q(a, s)$.
- Repeat until convergence.

Value iteration is **guaranteed to converge** to the optimal value function $V^*(s)$, and optimal action-value function $Q^*(a, s)$.

Deterministic Value Iteration

- If we know the solution to subproblems $V^*(s')$, then the solution $V^*(s)$ can be found by one-step lookahead

$$V^*(s) = \max_a \left[\mathbb{E}[r(a, s)] + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V^*(s') \right]$$

- The idea of value iteration is to apply these updates iteratively
- Start with final rewards and work backwards

Policy Iteration

Initialize $V(s)$ and a deterministic policy $\pi(s)$ for each state s . Suppose we know the reward function $r(a, s)$ and transition probabilities $P(s'|s, a)$.

- **Policy evaluation:** update the estimate of $V(s)$:

$$V^\pi(s) \leftarrow \mathbb{E}[r(\pi(s), s)] + \underbrace{\gamma \sum_{s' \in \mathcal{S}} P(s'|s, \pi(s)) V^\pi(s')}_{\text{Estimate of expected future reward}}$$

- **Policy improvement:** improve the policy at each state s :

$$\pi(s) \leftarrow \arg \max_a \left[\mathbb{E}[r(a, s)] + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V^\pi(s') \right]$$

- Repeat until π converges.
- Enjoys the same convergence rate as value iteration.

Q-learning

Q-Learning

What if we don't know $r(a, s)$ and $P(s'|s, a)$? Evolve the **action-value function** to find its optimal value $Q_{\pi^*}(a, s)$.

- Initialize our estimate of $Q(a, s)$ to some arbitrary value. We have dropped the dependence on π^* , as π^* is determined by $Q_{\pi^*}(a, s)$.
- After playing action a in state $s(t)$ and observing the next state $s(t+1)$, update Q as follows:

$$Q(a, s(t)) \leftarrow (1-\alpha) \underbrace{Q(a, s(t))}_{\text{old value}} + \alpha \underbrace{\left(r(a, s(t)) + \gamma \max_a Q(a, s(t+1)) \right)}_{\text{learned value}}.$$

Here α is the learning rate and $r(a, s(t))$ is our observed reward.

The term $\max_a Q(a, s(t+1))$ is our estimate of the expected future reward for the optimal policy π^* .

Q-learning has many variants: for instance, deep Q-learning uses a neural network to approximate Q .

Comparison with the Bellman Operator

$$\mathcal{T}(Q)(a, s) = \mathbb{E}[r(a, s)] + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) \max_a Q(a, s')$$

Since we no longer have access to the model $P(s'|s, a)$ and $r(a, s)$, we are replacing them by the observations, specifically

- replace $\mathbb{E}[r(a, s)]$ by the observed reward;
- replace $\sum_{s' \in \mathcal{S}} P(s'|s, a)$ by the observed state transition s'

In addition, make a conservative update instead of a full Bellman mapping to account for stochasticity:

$$Q(a, s) \leftarrow (1 - \alpha) \underbrace{Q(a, s)}_{\text{old value}} + \alpha \underbrace{\left(r(a, s(t)) + \gamma \max_a Q(a, s') \right)}_{\text{learned value}}.$$

Exploration vs. Exploitation

Given $Q(a, s)$, how do we choose our action a ?

- **Exploitation:** Take action $a^* = \arg \max_a Q(a, s)$. Given our current estimate of Q , we want to take what we think is the optimal action.
- **Exploration:** But we might not have a good estimate of Q , and we don't want to bias our estimate towards an action that turns out not to be optimal.
- **ϵ -Greedy:** With probability $1 - \epsilon$, choose $a^* = \arg \max_a Q(a, s)$, and otherwise choose a randomly. Usually, we decrease ϵ over time as additional exploration becomes less important.

Q-Learning Algorithm

Initialize our estimate of $Q(a, s)$ to some arbitrary value.

- Choose an action a using the ϵ -greedy policy.
- Observe reward $r(a, s(t))$ and state $s(t)$.
- Update Q as follows:

$$Q(a, s(t)) \leftarrow (1-\alpha) \underbrace{Q(a, s(t))}_{\text{old value}} + \alpha \underbrace{\left(r(a, s(t)) + \gamma \max_a Q(a, s(t+1)) \right)}_{\text{learned value}}.$$

- Repeat for T iterations.

Q-learning always learns an optimal policy, regardless of how a is chosen (you don't need to use ϵ -greedy)! Thus, it is an **off-policy** method.

Types of Machine Learning

Supervised Learning

- Training data: (x, y) (features, label) samples. We want to predict y to minimize a loss function.
- Regression, classification

Unsupervised Learning

- Training data: x (features) samples only. We want to find “similar” points in the x space.
- Clustering, PCA

Online/Sequential Learning

- Training data: (s, a, r) (state, action, reward) samples. We want to find the best sequence of decisions so as to maximize long-term reward.
- multi-armed bandits, reinforcement learning

You should know:

- Bellman optimality equation
- Value iteration, policy iteration, Q-learning

Course Summary

Machine learning is: the study of methods that

improve their performance

on some task

with experience


Machine Learning Pipeline



Examples

Example 1: Regression


How much should you sell your house for?



input: houses & features **learn:** $x \rightarrow y$ relationship **predict:** y (*continuous*)

Example 1: Regression

How much should you sell your house for?



input: houses & features **learn:** $x \rightarrow y$ relationship **predict:** y (*continuous*)

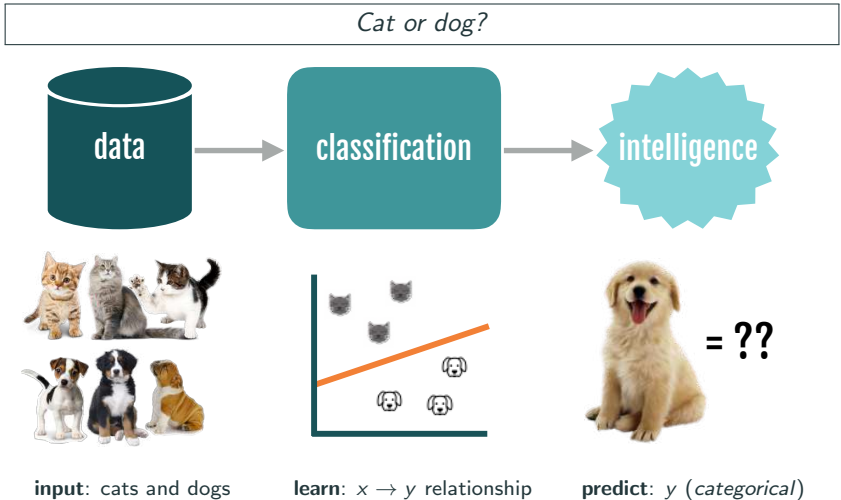
Models

- Linear regression
- Nonlinear models: neural networks/deep learning, decision trees, nearest neighbors

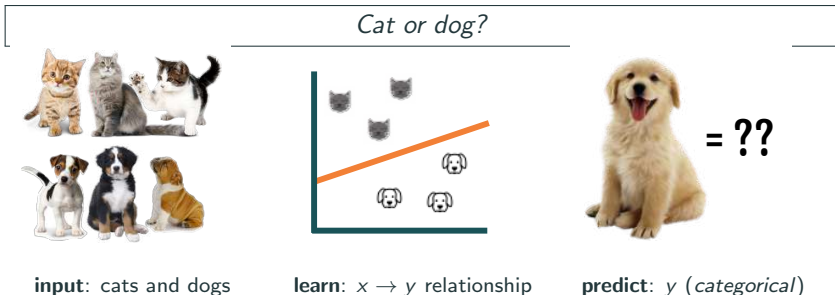
Methods

- MLE/MAP
- Gradient descent
- Ensemble methods (bagging, boosting)

Example 2: Classification



Example 2: Classification



Models

- Linear classification: Naïve Bayes, logistic regression, SVM
- Nonlinear models: kernel SVM, neural networks, decision trees, nearest neighbors

Methods


- MLE/MAP
- Gradient descent
- Ensemble methods (bagging, boosting)

Example 3: Clustering

How to segment an image?



input: raw pixels $\{x\}$



separate: $\{x\}$ into sets



output: cluster labels $\{z\}$

Example 3: Clustering

How to segment an image?



Models


- Clustering
- Gaussian mixture models

Methods


- k -means
- EM

Example 4: Embedding

How to reduce size of dataset?



input: large dataset $\{x\}$




find: sources of variation




return: representation $\{z\}$

Example 4: Embedding

How to reduce size of dataset?



input: large dataset $\{x\}$



find: sources of variation




return: representation $\{z\}$

Methods and Concepts

- PCA

Example 5: Sequential and Online Learning

How to learn sequentially in unknown environments?



Methods and Concepts

- Markov decision processes for reinforcement learning

Key Topics

Models

- Linear regression
- Linear classification:
Naïve Bayes,
logistic
regression, SVM
- Nonlinear
models: kernels,
neural networks
& deep learning,
decision trees
- Nearest
neighbors
- Clustering, GMM

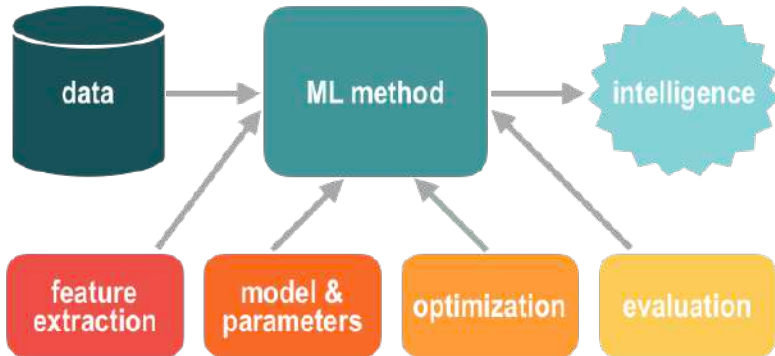
Methods

- (Stochastic)
gradient descent
- Boosting
- *k*-means
- EM
- PCA
- Distributed
learning

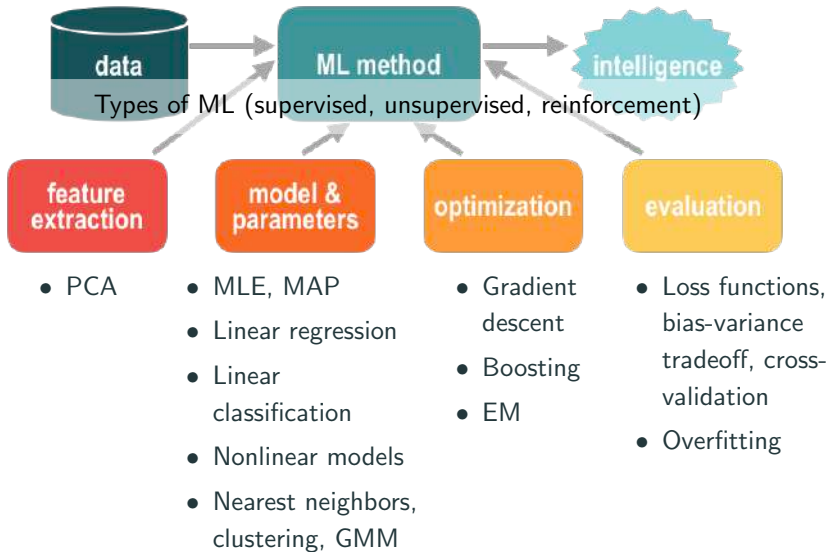
Concepts

- MLE, MAP
- Loss functions,
bias-variance
tradeoff,
cross-validation
- Bandits
- MDP, Bellman
equation
- Types of ML
(supervised,
unsupervised,
reinforcement)

Goal: Learn about ML Pipeline



Fitting the Course into the Pipeline



Please Complete Course Evaluations!

We will appreciate your constructive feedback on:

- Course Content/Structure
- Level of Math
- Homework and Pytorch assignment
- Our teaching
- and other suggestions/feedbacks

Completing course evaluations will earn you **1 bonus point towards your final grade**. Please upload the completion screenshot or email on Gradescope. Your answers will remain anonymous.

Thank you for taking the class!!

We hope you've enjoyed it as much as we have.

Final Exam Logistics

- **Final Exam** Friday of next week, May 2, 1 pm-4 pm ET (10am-1pm PT, 7pm-10pm CAT) in HOA 160 (Pitt), CMR F205 (Kigali), B23 118 (SV). **This is not the usual Pittsburgh lecture room.**
- All topics are included, with a focus on post-midterm topics
- Mix of true/false, multiple choice, and descriptive questions. No coding questions.
- 2 double-sided, US-letter or A4 sized handwritten cheat sheets allowed
- Calculators not permitted and not useful.
- Take the practice final exam, which will be posted on Piazza soon.

Office hours will be held as normal with exceptions noted on Piazza.

Recitation will also give you a chance to practice for the exam, and ask questions about the material.

Good luck!