

# **18-661 Introduction to Machine Learning**

## Neural Networks-III

---

Spring 2025

ECE – Carnegie Mellon University

1. Language Models and RNNs
2. Transformer Language Models
3. Stochastic Gradient Descent Convergence

# Language Models and RNNs

---

# What is Generative AI?

- So far, we have considered supervised learning tasks where we are given a training dataset of feature-label pairs  $(\mathbf{x}_n, y_n)$ , for  $n = 1, \dots, N$ . Our goal is to learn a function  $f(\mathbf{x}_n) \approx y_n$  that maps features to targets/labels
- In generative AI, we do not have explicit labels. Given a sequence of inputs  $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_t$  our goal is to predict the next element of the sequence  $\mathbf{x}_{t+1}$ .



# What is Generative AI?

- So far, we have considered supervised learning tasks where we are given a training dataset of feature-label pairs  $(\mathbf{x}_n, y_n)$ , for  $n = 1, \dots, N$ . Our goal is to learn a function  $f(\mathbf{x}_n) \approx y_n$  that maps features to targets/labels
- In generative AI, we do not have explicit labels. Given a sequence of inputs  $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_t$  our goal is to predict the next element of the sequence  $\mathbf{x}_{t+1}$ .
- Examples:
  - Next word prediction
  - Text generation given a prompt
  - Machine translation
  - Image/video generation from a description

# How does Generative AI work?

- We model the conditional distribution of the next token given the previous tokens:

$$\Pr(\mathbf{x}_{t+1} | \mathbf{x}_t, \dots \mathbf{x}_1)$$

using a neural network such as an RNN or transformer

- Then we sample from this probability distribution to generate  $\mathbf{x}_{t+1}$






Figure source: NVIDIA technical blog

# Recurrent Neural Networks (RNNs)



- Precursors to transformers, RNNs were widely used to model **temporal or sequential data** (e.g., natural language).
- Sequence of hidden states  $\mathbf{h}_t$  that depend on the current input  $\mathbf{x}_t$  and the previous hidden state  $\mathbf{h}_{t-1}$
- Output computation:  $\mathbf{o}_t = \psi(\mathbf{W}\mathbf{h}_t + b)$
- Hidden state computation:  $\mathbf{h}_t = \phi(\mathbf{V}\mathbf{h}_{t-1} + \mathbf{U}\mathbf{x}_t + c)$
- The weight matrices  $\mathbf{W}$ ,  $\mathbf{V}$ ,  $\mathbf{U}$  and biases  $b$  and  $c$  are trained using backpropagation on a dataset of sequences of varying lengths
- The predicted output  $\mathbf{o}_t$  becomes the next input  $\mathbf{x}_{t+1}$


# RNNs and Forgetting



- RNNs tend to forget information as they progress forward through the sequence
- This is due to weak or vanishing gradients as we move longer distance through the model
- For a long sentence where the beginning of the sentence has information about the subject, such forgetfulness can be catastrophic



# Long and Short-term Memory (LSTM)




- They combat the RNN forgetting issue via gates that decide whether to remember or forget information about the hidden states.
- But they still have drawbacks such as:
  - Difficulty with long-range dependencies (albeit less than RNNs)
  - Even though they solve the vanishing gradient problem, they suffer from exploding gradients
  - Inherently serial computation makes it harder to train

# Transformer Language Models

---


# Transformers: Attention Mechanism



- RNNs and LSTM maintain a fixed length hidden state to represent the history of a sequence.
- Instead, the attention mechanism in Transformers looks at all previous token when predicting the next token
- The 'attention' that it pays to each token is computed using the attention mechanism

# Transformers

- **Encoder-decoder** architecture: learn a representation of each input in the sequence, then decode to predict next entry in the output sequence
  - Autoregressive structure: takes previously generated outputs as inputs
  - Attach an attention weight to each entry of each input representation
  - Self-attention between each layer
- Much larger and slower to train, but usually gives good performance



# Large Language Models

- “Generative pre-trained transformer” models: *generate* language outputs based on *pre-training* of *transformer*-based architectures on a massive corpus of language data
- Classification task: next-word prediction (run many times)
  - Tokenization: divide text into ‘tokens’ of similar length/information
  - Predict the next token based on the preceding sequence of tokens (typically 1M long)
- Self-/semi-supervised models: generate supervisory signals (“labels”) based on output of currently trained model
- Other generative models can generate images, videos, etc.  
Multi-modal models can, e.g., use a text input to generate an image.

# Pre-Training and Finetuning

- Modern deep learning models are **too expensive to train from scratch** (GPT-4 likely cost millions of dollars to train!) As an example, Llama-7B, 13-B, etc. have billions of parameters.
- Pre-trained **foundation models** capture essential patterns and can be finetuned to specific datasets
  - Types of language, e.g., coding tools or translation tasks
  - Types of images, e.g., generating images of a certain style
- Foundation models can be trained further on a new dataset
  - Layer freezing or prompt engineering
  - “Warm start” initialization to the usual SGD-based training steps
- **Prune, quantize, or compress** foundation models to fit them or train them on smaller devices.

1. Language Models and RNNs
2. Transformer Language Models
3. Stochastic Gradient Descent Convergence

# Stochastic Gradient Descent Convergence

---



# SGD is at the core of Machine Learning!

We use it to train the model parameters  $\mathbf{w}$  in

- Linear Regression:  $y = \mathbf{w}^\top \mathbf{x}$
- Logistic Regression:  $y = \sigma(\mathbf{w}^\top \mathbf{x})$
- Neural Networks:  $y = NN(\mathbf{x}; \mathbf{w})$

For each problem, we define a loss function  $F(\mathbf{w})$  to measure the error in the predicted output, and then update  $\mathbf{w}$  according to:

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \eta g(\mathbf{w})$$

# SGD is at the core of Machine Learning!

For each problem, we define a loss function  $F(\mathbf{w})$  to measure the error in the predicted output, and then update  $\mathbf{w}$  according to:


$$\mathbf{w}_{t+1} = \mathbf{w}_t - \eta g(\mathbf{w})$$

The gradient  $g(\mathbf{w})$  can be


- Full gradient,  $\nabla F(\mathbf{w}_t) = \frac{1}{N} \sum_{i=1}^N \nabla f(\mathbf{w}_t; \xi)$  computed over the whole dataset
- Stochastic gradient  $\nabla f(\mathbf{w}_t; \xi)$  for a randomly chosen sample  $\xi$
- Mini-batch stochastic gradient  $g(\mathbf{w}; \xi) = \frac{1}{b} \sum_{i=1}^b \nabla f(\mathbf{w}; \xi_i)$ , computed using a batch  $\xi$  of  $b$  samples chosen at random

Let us analyze the time SGD takes to reach an  $\epsilon$  error


# A $c$ -strongly Convex and $L$ -Smooth Function



—  $f : \mathbb{R}^d \rightarrow \mathbb{R}$   
CONVEX FUNCTION



—  $g : \mathbb{R}^d \rightarrow \mathbb{R}$   
STRONGLY CONVEX  
FUNCTION



—  $g : \mathbb{R}^d \rightarrow \mathbb{R}$   
STRONGLY SMOOTH  
CONVEX FUNCTION

Satisfies the upper and lower bounds given by

$$F(\mathbf{w}) \leq F(\mathbf{y}) + \nabla F(\mathbf{y})^\top (\mathbf{w} - \mathbf{y}) + \frac{L}{2} \|\mathbf{w} - \mathbf{y}\|^2$$

$$F(\mathbf{w}) \geq F(\mathbf{y}) + \nabla F(\mathbf{y})^\top (\mathbf{w} - \mathbf{y}) + \frac{1}{2} c \|\mathbf{w} - \mathbf{y}\|^2 \text{ for all } \mathbf{w}, \mathbf{y} \in \mathbb{R}^d$$

# Convergence Analysis of GD

## Convergence of GD

For a  $c$ -strongly convex and  $L$ -smooth function, if the learning rate  $\eta < \frac{1}{L}$  and the starting point is  $\mathbf{w}_0$  then  $F(\mathbf{w}_t)$  after  $t$  gradient descent iterations is bounded as

$$F(\mathbf{w}_t) - F(\mathbf{w}^*) \leq (1 - \eta c)^t (F(\mathbf{w}_0) - F(\mathbf{w}^*))$$

How many iterations do we need to converge to reach error  $F(\mathbf{w}_t) - F(\mathbf{w}^*) = \epsilon$ ?

$$(1 - \eta c)^t (F(\mathbf{w}_0) - F(\mathbf{w}^*)) \leq \epsilon$$

$$t \log(1 - \eta c) + \log(F(\mathbf{w}_0) - F(\mathbf{w}^*)) \leq \log(\epsilon)$$

$$t \log(1/(1 - \eta c)) - \log(F(\mathbf{w}_0) - F(\mathbf{w}^*)) \geq \log\left(\frac{1}{\epsilon}\right)$$

$$t = O\left(\log\left(\frac{1}{\epsilon}\right)\right)$$

# Convergence Analysis of GD

## Convergence of GD

For a  $c$ -strongly convex and  $L$ -smooth function, if the learning rate  $\eta < \frac{1}{L}$  and the starting point is  $\mathbf{w}_0$  then  $F(\mathbf{w}_t)$  after  $t$  gradient descent iterations is bounded as

$$F(\mathbf{w}_t) - F(\mathbf{w}^*) \leq (1 - \eta c)^t (F(\mathbf{w}_0) - F(\mathbf{w}^*))$$

How does the convergence speed depend on

- Learning rate  $\eta$  – Converges faster for larger  $\eta$  as long as  $\eta < \frac{1}{L}$
- Lipschitz smoothness  $L$  – Converges faster for smaller  $L$  because we can set a higher  $\eta$
- Strong convexity parameter  $c$  – Converges faster for larger  $c$

# Convergence Analysis of Mini-batch SGD

# Assumptions on the Stochastic Gradients

Since we are using noisy gradients, we need the following assumptions on them

- **Unbiased Gradients:** The stochastic gradient  $\nabla f(\mathbf{w}; \xi)$  is an unbiased estimate of  $\nabla F(\mathbf{w})$ , that is,

$$\mathbb{E}_{\xi}[\nabla f(\mathbf{w}; \xi)] = \nabla F(\mathbf{w})$$

- **Bounded Variance:** The stochastic gradient  $\nabla f(\mathbf{w}; \xi)$  has bounded variance, that is,

$$\text{Var}(\nabla f(\mathbf{w}; \xi)) \leq \sigma^2$$

which implies that the variance of a mini-batch gradient is:

$$\text{Var}(g(\mathbf{w}; \xi)) \leq \frac{\sigma^2}{b}$$

# Convergence Analysis of Mini-batch SGD

## Convergence of Mini-batch SGD

For a  $c$ -strongly convex and  $L$ -smooth function, if the learning rate  $\eta < \frac{1}{L}$  and the starting point is  $\mathbf{w}_0$  then  $F(\mathbf{w}_t)$  after  $t$  gradient descent iterations is bounded as

$$\mathbb{E}[F(\mathbf{w}_t)] - F(\mathbf{w}^*) - \frac{\eta L \sigma^2}{2cb} \leq (1 - \eta c)^t \left( \mathbb{E}[F(\mathbf{w}_0)] - F(\mathbf{w}^*) - \frac{\eta L \sigma^2}{2cb} \right)$$

- For batch GD, as  $t \rightarrow \infty$ , the objective  $F(\mathbf{w}_t) \rightarrow F(\mathbf{w}^*)$
- For mini-batch SGD, as  $t \rightarrow \infty$ , we will be left with an error floor  $\mathbb{E}[F(\mathbf{w}_t)] - F(\mathbf{w}^*) \rightarrow \frac{\eta L \sigma^2}{2cb}$ .
- This is the price that we pay for noisy gradients, that is the variance bound being  $\sigma^2 \geq 0$



# Effect of Mini-batch Size on the Error Floor

## Convergence of Mini-batch SGD

For a  $c$ -strongly convex and  $L$ -smooth function, if the learning rate  $\eta < \frac{1}{L}$  and the starting point is  $\mathbf{w}_0$  then  $F(\mathbf{w}_t)$  after  $t$  gradient descent iterations is bounded as

$$\mathbb{E}[F(\mathbf{w}_t)] - F(\mathbf{w}^*) - \frac{\eta L \sigma^2}{2cb} \leq (1 - \eta c)^t \left( \mathbb{E}[F(\mathbf{w}_0)] - F(\mathbf{w}^*) - \frac{\eta L \sigma^2}{2bc} \right)$$

- Recall that

$$g(\mathbf{w}, \xi) = \frac{1}{b} \sum_{n \in S} \nabla f(\mathbf{w}),$$

- And the bounded variance assumption is  $\text{Var}(g(\mathbf{w}; \xi)) \leq \sigma^2/b$
- When we increase the mini-batch size  $b$ , the error floor  $\frac{\eta L \sigma^2}{2cb}$  reduces.

# Effect of Learning Rate on Convergence Speed and Error Floor

## Convergence of Mini-batch SGD

For a  $c$ -strongly convex and  $L$ -smooth function, if the learning rate  $\eta < \frac{1}{L}$  and the starting point is  $\mathbf{w}_0$  then  $F(\mathbf{w}_t)$  after  $t$  gradient descent iterations is bounded as

$$\mathbb{E}[F(\mathbf{w}_t)] - F(\mathbf{w}^*) - \frac{\eta L \sigma^2}{2cb} \leq (1 - \eta c)^t \left( \mathbb{E}[F(\mathbf{w}_0)] - F(\mathbf{w}^*) - \frac{\eta L \sigma^2}{2cb} \right)$$

- The convergence speed, represents by  $(1 - \eta c)$  only depends on the strong convexity parameter  $c$  and learning rate  $\eta$ , not on the mini-batch size  $b$
- As  $\eta$  increases, the algorithm converges faster
- But, as  $\eta$  increases, the error floor  $\frac{\eta L \sigma^2}{2cb}$  also increases


# How do we achieve zero error floor?

## Convergence of Mini-batch SGD

For a  $c$ -strongly convex and  $L$ -smooth function, if the learning rate  $\eta < \frac{1}{L}$  and the starting point is  $\mathbf{w}_0$  then  $F(\mathbf{w}_t)$  after  $t$  mini-batch SGD iterations is bounded as

$$\mathbb{E}[F(\mathbf{w}_t)] - F(\mathbf{w}^*) - \frac{\eta L \sigma^2}{2cb} \leq (1 - \eta c)^t \left( \mathbb{E}[F(\mathbf{w}_0)] - F(\mathbf{w}^*) - \frac{\eta L \sigma^2}{2cb} \right)$$

**KEY IDEA:** Decay the learning rate  $\eta$  (denoted by  $\alpha$  in the figure below) by 2 whenever  $F$  is 2 times its error floor  $\frac{\eta L \sigma^2}{2cb}$ .



# How do we achieve zero error floor?

## Convergence of Mini-batch SGD

For a  $c$ -strongly convex and  $L$ -smooth function, if the learning rate  $\eta < \frac{1}{L}$  and the starting point is  $\mathbf{w}_0$  then  $F(\mathbf{w}_t)$  after  $t$  mini-batch SGD iterations is bounded as

$$\mathbb{E}[F(\mathbf{w}_t)] - F(\mathbf{w}^*) - \frac{\eta L \sigma^2}{2cb} \leq (1 - \eta c)^t \left( \mathbb{E}[F(\mathbf{w}_0)] - F(\mathbf{w}^*) - \frac{\eta L \sigma^2}{2cb} \right)$$

**KEY IDEA:** Decay the learning rate  $\eta$

- If  $\eta = \eta_0/t$ , then we can show that  $\mathbb{E}[F(\mathbf{w}_t)] - F(\mathbf{w}^*) \leq O(1/t)$
- Thus, the number of iterations required an error  $\epsilon$  is  $O(1/\epsilon)$
- In contrast, with GD we need only  $O(\log(1/\epsilon))$  iterations to reach an  $\epsilon$  error

You should know:

- Language models and RNNs
- Transformer language models