

# 18-661 Introduction to Machine Learning

Pytorch

---

Arian Raje

Spring 2025

ECE – Carnegie Mellon University

1. A Small History Lesson
2. What is Pytorch?
3. Pytorch Examples

**Disclaimer:** this lecture will not appear on your final exam, though some content, in particular PyTorch, may be used on the Homework

# Some History About Neural Networks

---

# A Grossly Abbreviated History of AI

## 4 Time Periods of Artificial Intelligence/Neural Networks -

### 1. Early Stages (1950s-1980s)

- Turing Test
- Perceptron

### 2. AI Winter (1980s-2000s)

- *Perceptrons: An Introduction to Computational Geometry* by Marvin Minsky

### 3. Resurgence and GPU Acceleration (2000s-2010s)

- DARPA
- IBM Watson
- CUDA

### 4. Modern AI (2010s-2020s)

- Pytorch/Tensorflow
- AlexNet
- ChatGPT

# A Grossly Abbreviated History of AI

## 4 Time Periods of Artificial Intelligence/Neural Networks -

### 1. Early Stages (1950s-1980s)

- Turing Test
- Perceptron

### 2. AI Winter (1980s-2000s)

- *Perceptrons: An Introduction to Computational Geometry* by Marvin Minsky

### 3. Resurgence and GPU Acceleration (2000s-2010s)

- DARPA
- IBM Watson
- CUDA

### 4. Modern AI (2010s-2020s)

- PyTorch / TensorFlow
- AlexNet
- ChatGPT

# What is Pytorch?

---

# What is PyTorch?

- **Open-Source Machine Learning Library** developed by Facebook's AI Research (FAIR) team
- Built primarily for **Deep Learning** tasks (vision, NLP, recommender systems, etc.)
- Provides:
  - **Automatic differentiation**
  - **GPU acceleration**
- Pythonic, easy to debug, and flexible

# Automatic Differentiation (Autodiff)

- **Idea:**

- Computes exact derivatives by “tracing” operations during the forward pass.
- Replaces manual/finite-difference methods (less error-prone, more efficient).

- **Key Benefits:**

- No manual math or derivative coding.
- Works seamlessly with complex, dynamic architectures.

- **Typical Workflow (e.g., PyTorch):**

1. Define forward pass (model + loss).
2. Operations are recorded in a graph.
3. Call `.backward()` for gradients w.r.t. parameters.



# Example w/ Sequential Model

```
class MySequentialModel(nn.Module):
    def __init__(self):
        super(MySequentialModel, self).__init__()
        self.layers = nn.Sequential(
            nn.Linear(10, 5),
            nn.ReLU(),
            nn.Linear(5, 2)
        )

    def forward(self, x):
        # Forward pass: simply feed x through the Sequential layers
        return self.layers(x)
```

MySequentialModel definition

```
# 3. Perform a forward pass
output = model(x)

# 4. Define a dummy target and loss function
target = torch.randn(3, 2) # same shape as output
criterion = nn.MSELoss()

# 5. Calculate the loss
loss = criterion(output, target)

# 6. Backward pass: compute gradients w.r.t. all parameters
loss.backward()
```

Forward pass, loss, and backward call

# Key Benefit of Autodiff

You don't need to have an **explicit** backward function. Only defining a forward function suffices.

- **Why GPUs?**

- Optimized for parallel operations (fast matrix math)
- Huge speedups for neural network training

- **Using GPUs in PyTorch:**

- `device = torch.device("cuda" if torch.cuda.is_available() else "cpu")`
- `model.to(device)` and `x.to(device)` before training
- Automatic differentiation runs on GPU

- **Multi-GPU / Distributed:**

- `nn.DataParallel` or `DistributedDataParallel`
- Scale across multiple GPUs/machines

# GPU vs. CPU


	<b>CPU</b>	<b>GPU</b>
<b>Core Count</b>	Few (4–32)	Hundreds to Thousands
<b>Parallelism</b>	Limited	Massively Parallel
<b>Memory Bandwidth</b>	Lower	Higher
<b>Primary Use</b>	General-purpose	Graphics / ML Acceleration

**Question:** Why wouldn't you always use GPUs?

The cores in CPUs are capable of performing very different instructions. GPU cores have to perform the same instruction.

# A Typical Deep Learning Pipeline

- The CPU is usually used for data preprocessing only.
- All parameter and gradient computations take place on the GPU
- Data loading and preprocessing should be pipelined to avoid impacting runtime.



# Advantages of Pytorch-like Frameworks

## 2 Key Advantages -

- Automatic differentiation
- GPU acceleration

Previously, you would have to write custom CUDA code and calculate gradients by hand. With these frameworks, all of that is done for you.

## Similar Frameworks -

- Theano → TensorFlow
- Torch → Pytorch
- JAX



## Other Related Frameworks

Framework	Developer	Why Use It?
<b>TensorFlow</b>	Google Brain	Large ecosystem; production-ready (TensorFlow Serving / Lite)
<b>PyTorch</b>	Meta AI (Facebook)	Pythonic and dynamic; strong research and industry adoption
<b>JAX</b>	Google	High-performance autodiff via just-in-time compilation (XLA)

# PyTorch Example

---