

# **18-661 Introduction to Machine Learning**

## Support Vector Machines (SVM) – I

---

Spring 2025

ECE – Carnegie Mellon University

# Announcements and Reminders


- Homework 2 is posted and due February 21.
- No recitation on Friday this week.
- Midterm exam in two weeks, on February 26. More details to be announced in Wednesday's lecture next week.
  - Similar format to the mini-exam, but longer (110 minutes).
  - Covers all topics through the February 19 lecture.

1. Review of Non-linear Classification Boundaries
2. Review of Multi-class Logistic Regression
3. Evaluating Classification Methods
4. Support Vector Machines (SVM): Intuition
5. SVM: Max-Margin Formulation
6. SVM: Hinge Loss Formulation

# **Review of Non-linear Classification Boundaries**

---

# How to Handle More Complex Decision Boundaries?




- This data is not linearly separable...
- Use **non-linear basis functions** to add more features (and hope the data is separable in the augmented feature space).

# Solution to Overfitting: Regularization

- Add regularization term to be cross entropy loss function

$$\mathcal{E}(\mathbf{w}) = - \sum_n \{y_n \log \sigma(\mathbf{w}^\top \mathbf{x}_n) + (1-y_n) \log [1 - \sigma(\mathbf{w}^\top \mathbf{x}_n)]\} + \underbrace{\frac{1}{2} \lambda \|\mathbf{w}\|_2^2}_{\text{regularization}}$$

- Perform gradient descent on this regularized function.
- Often, we do **NOT** regularize the bias term  $w_0$ .




# Review of Multi-class Logistic Regression

---

# Three Approaches


- One-versus-all
- One-versus-one
- Multinomial regression





# The One-versus-Rest or One-versus-All Approach

- For each class  $c$ , change the problem into binary classification
  1. Relabel training data with label  $c$ , into POSITIVE (or '1').
  2. Relabel all the rest data into NEGATIVE (or '0').
- Repeat this multiple times: Train  $C$  binary classifiers, using logistic regression to differentiate the two classes each time.
- There is ambiguity in some of the regions (the 4 triangular areas)...




# The One-versus-Rest or One-versus-All Approach

How to combine these linear decision boundaries?


- Use the **confidence estimates**  $\Pr(y = 1|\mathbf{x}) = \sigma(\mathbf{w}_1^\top \mathbf{x})$ ,  
...  $\Pr(y = C|\mathbf{x}) = \sigma(\mathbf{w}_C^\top \mathbf{x})$
- Declare class  $c^*$  that maximizes

$$c^* = \arg \max_{c=1,\dots,C} \Pr(y = c|\mathbf{x}) = \sigma(\mathbf{w}_c^\top \mathbf{x})$$




# The One-versus-One Approach

- For each **pair** of classes  $c$  and  $c'$ , change the problem into binary classification.
  1. Relabel training data with label  $c$ , into POSITIVE (or '1')
  2. Relabel training data with label  $c'$  into NEGATIVE (or '0')
  3. **Disregard** all other data



# The One-versus-One Approach

- How many binary classifiers for  $C$  classes?  $C(C - 1)/2$
- How to combine their outputs?
- Given  $\mathbf{x}$ , count the  $C(C - 1)/2$  votes from outputs of all binary classifiers and declare the winner as the predicted class.
- Use confidence scores to resolve ties.



# Multinomial Logistic Regression

- **Model:** For each class  $c$ , we have a parameter vector  $\mathbf{w}_c$  and model the posterior probability as:

$$P(c|\mathbf{x}) = \frac{e^{\mathbf{w}_c^\top \mathbf{x}}}{\sum_{c'} e^{\mathbf{w}_{c'}^\top \mathbf{x}}} \quad \leftarrow \quad \text{This is called the } \textit{softmax} \text{ function.}$$

- **Decision boundary:** Assign  $\mathbf{x}$  with the label that is the maximum of posterior:

$$\arg \max_c P(c|\mathbf{x}) \rightarrow \arg \max_c \mathbf{w}_c^\top \mathbf{x}.$$

# Parameter Estimation for Multinomial Logistic Regression

**Discriminative approach:** Maximize conditional likelihood

$$\log P(\mathcal{D}) = \sum_n \log P(y_n | \mathbf{x}_n)$$

We will change  $y_n$  to  $\mathbf{y}_n = [y_{n1} \ y_{n2} \ \cdots \ y_{nC}]^\top$ , a  $C$ -dimensional vector using 1-of- $C$  encoding.

$$y_{nc} = \begin{cases} 1 & \text{if } y_n = c \\ 0 & \text{otherwise} \end{cases}$$

Ex: if  $y_n = 2$ , then,  $\mathbf{y}_n = [0 \ \mathbf{1} \ 0 \ 0 \ \cdots \ 0]^\top$ .

$$\Rightarrow \sum_n \log P(y_n | \mathbf{x}_n) = \sum_n \log \prod_{c=1}^C P(c | \mathbf{x}_n)^{y_{nc}} = \sum_n \sum_c y_{nc} \log P(c | \mathbf{x}_n)$$

# Cross-entropy Error Function

**Definition:** negative log-likelihood

$$\begin{aligned}\mathcal{E}(\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_C) &= - \sum_n \sum_c y_{nc} \log P(c|\mathbf{x}_n) \\ &= - \sum_n \sum_c y_{nc} \log \left( \frac{e^{\mathbf{w}_c^\top \mathbf{x}_n}}{\sum_{c'} e^{\mathbf{w}_{c'}^\top \mathbf{x}_n}} \right)\end{aligned}$$

## Properties of cross-entropy

- Convex in the  $\mathbf{w}_c$  vectors, therefore unique global optimum
- Optimization requires numerical procedures, analogous to those used for binary logistic regression.

# Evaluating Classification Methods

---



$$\mathcal{E}(\mathbf{w}) = - \sum_n \{y_n \log \sigma(\mathbf{w}^\top \mathbf{x}_n) + (1 - y_n) \log[1 - \sigma(\mathbf{w}^\top \mathbf{x}_n)]\}$$

- Easy to optimize!
- Average loss over the (training, validation, test) dataset
- ...but what does it mean?

# Interpretable Classification Metrics

True positive	False positive
False negative	True negative

- Measure the accuracy within each class
- Accounts for imbalance between classes

These metrics are **difficult to optimize directly**, but they have the advantage of being easily interpretable.

- **Sensitivity**: true positive rate

$$\text{TPR} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$


- **Specificity**: true negative rate

$$\text{TNR} = \frac{\text{TN}}{\text{TN} + \text{FP}}$$

- **Precision**: positive predictive value

$$\text{PPV} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

# Combining These Metrics: the ROC Curve



Receiver Operating Characteristic  
(ROC)

- Define a “threshold” for the positive/negative split
- Increasing the threshold: more samples are predicted to be positive
- **Area Under the ROC Curve:** want this as large as possible

1. Review of Non-linear Classification Boundaries
2. Review of Multi-class Logistic Regression
3. Evaluating Classification Methods
4. Support Vector Machines (SVM): Intuition
5. SVM: Max-Margin Formulation
6. SVM: Hinge Loss Formulation

# Support Vector Machines (SVM): Intuition

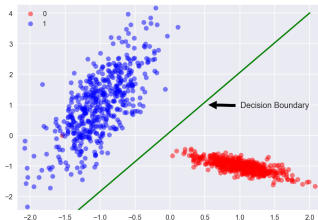

---

# Why Do We Need SVM?

Alternative to Logistic Regression and Naïve Bayes.

- Logistic regression and Naïve Bayes train over the whole dataset.
- These can require a lot of memory in high-dimensional settings.
- SVM can give a better and more efficient solution.
- SVM is one of the most powerful and commonly used ML algorithms.

# Binary Logistic Regression



- We only need to know if  $p(\mathbf{x}) > 0.5$  or  $< 0.5$ .
- We **don't** (always) need to know how far  $\mathbf{x}$  is from this boundary.

**How can we use this insight to improve the classification algorithm?**

- What if we just looked at the boundary?
- Maybe then we could ignore some of the samples?

# Advantages of SVM

We will see later that SVM:

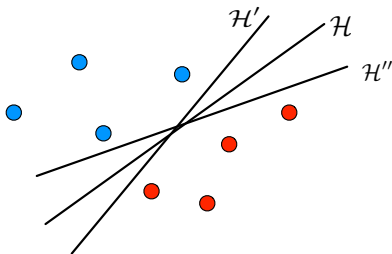
1. Maximizes distance of training points from the boundary
2. Only requires a subset of the training points.
3. Is less sensitive to outliers.
4. Scales better with high-dimensional data.
5. Generalizes well to many nonlinear models.



## **SVM: Max-Margin Formulation**

---


# Binary Classification: Finding a Linear Decision Boundary



- Input features  $\mathbf{x}$ .
- Decision boundary is a hyperplane  $\mathcal{H} : \mathbf{w}^\top \mathbf{x} + b = 0$ .

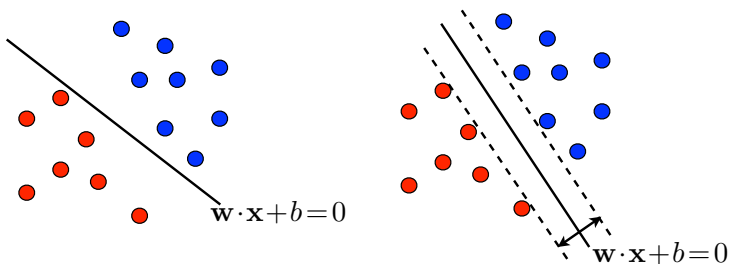
# Intuition: Where to Put the Decision Boundary?

- Consider a *separable* training dataset (e.g., with two features)
- There are an **infinite** number of decision boundaries  
 $\mathcal{H} : \mathbf{w}^\top \mathbf{x} + b = 0$ !



- Which one should we pick?

# Intuition: Where to Put the Decision Boundary?





Find a decision boundary in the '*middle*' of the two classes that:

- Perfectly classifies the training data
- Is as far away from every training point as possible

Let us apply this intuition to build a classifier that **maximizes the margin** between training points and the decision boundary.

# First, Some Vector Geometry

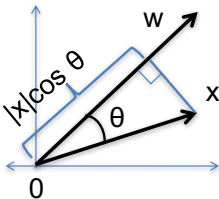
What is a hyperplane?



- General equation is  $\mathbf{w}^\top \mathbf{x} + b = 0$
- Divides the space in half, i.e.,  $\mathbf{w}^\top \mathbf{x} + b > 0$  and  $\mathbf{w}^\top \mathbf{x} + b < 0$
- A hyperplane is a line in 2D and a plane in 3D
- $\mathbf{w} \in \mathbb{R}^d$  is a non-zero normal vector


# Vector Norms and Inner Products

- Given two vectors  $\mathbf{w}$  and  $\mathbf{x}$ , what is their inner product?
- Inner Product  $\mathbf{w}^\top \mathbf{x} = w_1x_1 + w_2x_2 + \cdots + w_dx_d$



- Inner Product  $\mathbf{w}^\top \mathbf{x}$  is also equal to  $\|\mathbf{w}\| \|\mathbf{x}\| \cos \theta$
- $\mathbf{w}^\top \mathbf{w} = \|\mathbf{w}\|^2$
- If  $\mathbf{w}$  and  $\mathbf{x}$  are perpendicular, then  $\theta = \pi/2$ , and thus the inner product is zero.


# Normal Vector of a Hyperplane



**Vector  $w$  is normal to the hyperplane. Why?**

- If  $p$  and  $q$  are both on the line, then  $w^\top p + b = w^\top q + b = 0$ .
- Then  $w^\top (p - q) = w^\top p - w^\top q = -b - (-b) = 0$
- $p - q$  is an arbitrary vector parallel to the line, thus  $w$  is orthogonal
- $w^* = \frac{w}{\|w\|_2}$  is the unit normal vector

# Distance from a Hyperplane



## How to find the distance from $\mathbf{a}$ to the hyperplane?

- We want to find distance between  $\mathbf{a}$  and line in the direction of  $\mathbf{w}^*$ .
- If we define point  $\mathbf{a}_0$  on the line, then this distance corresponds to length of  $\mathbf{a} - \mathbf{a}_0$  in direction of  $\mathbf{w}^*$ , which equals  $\mathbf{w}^{*\top}(\mathbf{a} - \mathbf{a}_0)$ .
- We know  $\mathbf{w}^\top \mathbf{a}_0 = -b$  since  $\mathbf{w}^\top \mathbf{a}_0 + b = 0$ .
- Then the distance equals  $\frac{1}{\|\mathbf{w}\|_2}(\mathbf{w}^\top \mathbf{a} + b)$ .



# Distance from a Point to Decision Boundary

The *unsigned* distance from a point  $\mathbf{x}$  to the decision boundary (hyperplane)  $\mathcal{H}$  is

$$d_{\mathcal{H}}(\mathbf{x}) = \frac{|\mathbf{w}^{\top} \mathbf{x} + b|}{\|\mathbf{w}\|_2}$$

How to remove the absolute value  $|\cdot|$ ?

**Notation changes from Logistic Regression:** Use  $y = +1$  to represent positive label and  $y = -1$  for negative label.

Then, exploiting the fact that the decision boundary classifies every point in the training dataset correctly, we have  $(\mathbf{w}^{\top} \mathbf{x} + b)$  and  $\mathbf{x}$ 's label  $y$  must have the same sign. So we get

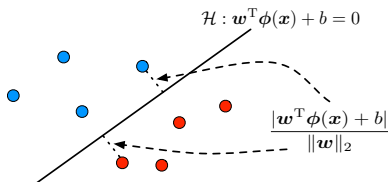
$$d_{\mathcal{H}}(\mathbf{x}) = \frac{y[\mathbf{w}^{\top} \mathbf{x} + b]}{\|\mathbf{w}\|_2}$$

# Defining the Margin

## Margin

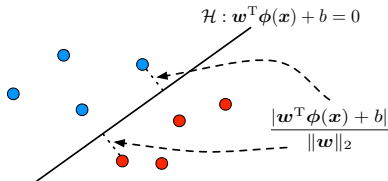
Smallest distance between the hyperplane and all training points

$$\text{MARGIN}(\mathbf{w}, b) = \min_n \frac{y_n[\mathbf{w}^\top \mathbf{x}_n + b]}{\|\mathbf{w}\|_2}$$



How can we use this to find the SVM solution?

# Optimizing the Margin



## How should we pick $(\mathbf{w}, b)$ based on its margin?

We want a decision boundary that is as far away from all training points as possible, so we to *maximize* the margin!

$$\max_{\mathbf{w}, b} \left( \min_n \frac{y_n [\mathbf{w}^T \mathbf{x}_n + b]}{\|\mathbf{w}\|_2} \right) = \max_{\mathbf{w}, b} \left( \frac{1}{\|\mathbf{w}\|_2} \min_n y_n [\mathbf{w}^T \mathbf{x}_n + b] \right)$$

Only involves points near the boundary (more on this later).

## Margin

Smallest distance between the hyperplane and all training points

$$\text{MARGIN}(\mathbf{w}, b) = \min_n \frac{y_n[\mathbf{w}^\top \mathbf{x}_n + b]}{\|\mathbf{w}\|_2}$$

Consider three hyperplanes

$$(\mathbf{w}, b) \quad (2\mathbf{w}, 2b) \quad (.5\mathbf{w}, .5b)$$

Which one has the largest margin?

- The MARGIN doesn't change if we scale  $(\mathbf{w}, b)$  by a constant  $c$
- $\mathbf{w}^\top \mathbf{x} + b = 0$  and  $(c\mathbf{w})^\top \mathbf{x} + (cb) = 0$ : same decision boundary!
- Can we further constrain the problem so as to get a unique solution  $(\mathbf{w}, b)$ ?

# Rescaled Margin

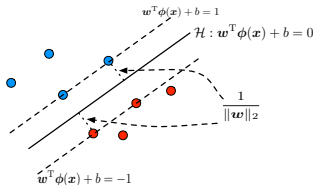
We can further constrain the problem by scaling  $(\mathbf{w}, b)$  such that

$$\min_n y_n [\mathbf{w}^\top \mathbf{x}_n + b] = 1.$$

Note that there always exists a scaling for which this is true. We've fixed the numerator in the  $\text{MARGIN}(\mathbf{w}, b)$  equation, and we have:

$$\text{MARGIN}(\mathbf{w}, b) = \frac{\min_n y_n [\mathbf{w}^\top \mathbf{x}_n + b]}{\|\mathbf{w}\|_2} = \frac{1}{\|\mathbf{w}\|_2}$$

Hence the points closest to the decision boundary are at distance  $\frac{1}{\|\mathbf{w}\|_2}$ .



# SVM: Max-margin Formulation for Separable Data

We thus want to solve:

$$\max_{\mathbf{w}, b} \underbrace{\frac{1}{\|\mathbf{w}\|_2}}_{\text{margin}} \quad \text{such that} \quad \underbrace{\min_n y_n [\mathbf{w}^\top \mathbf{x}_n + b]}_{\text{scaling of } \mathbf{w}, b} = 1$$

which is equivalent to

$$\max_{\mathbf{w}, b} \frac{1}{\|\mathbf{w}\|_2} \quad \text{such that} \quad y_n [\mathbf{w}^\top \mathbf{x}_n + b] \geq 1, \quad \forall n$$

This is further equivalent to

$$\begin{aligned} \min_{\mathbf{w}, b} \quad & \frac{1}{2} \|\mathbf{w}\|_2^2 \\ \text{s.t.} \quad & y_n [\mathbf{w}^\top \mathbf{x}_n + b] \geq 1, \quad \forall n \end{aligned}$$

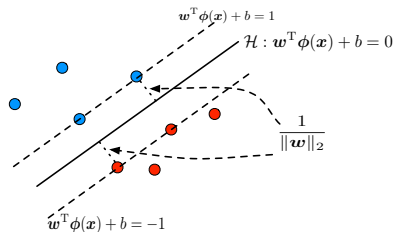
Given our geometric intuition, SVM is called a **max margin** (or large margin) classifier. The constraints are called **large margin constraints**.

# Support Vectors: A First Look

## SVM formulation for separable data

$$\min_{\mathbf{w}, b} \quad \frac{1}{2} \|\mathbf{w}\|_2^2$$

$$\text{s.t. } y_n[\mathbf{w}^\top \mathbf{x}_n + b] \geq 1, \quad \forall n$$



Two types of training data, based on the situations of the constraint:

- “=”:  $y_n[\mathbf{w}^\top \mathbf{x}_n + b] = 1$ . These training data points are called “support vectors”, which have the minimum distance  $(\frac{1}{\|\mathbf{w}\|})$  to the boundary.
- “>”:  $y_n[\mathbf{w}^\top \mathbf{x}_n + b] > 1$ . Distance to the boundary is larger than the minimum. Removing these data points does not affect the optimal solution (more on this next lecture).

# SVM for Non-separable Data

## SVM formulation for separable data

$$\begin{aligned} \min_{\mathbf{w}, b} \quad & \frac{1}{2} \|\mathbf{w}\|_2^2 \\ \text{s.t.} \quad & y_n[\mathbf{w}^\top \mathbf{x}_n + b] \geq 1, \quad \forall n \end{aligned}$$

## Non-separable setting

In practice our training data may not be separable. What issues arise with the optimization problem above when data is not separable?

- For every  $\mathbf{w}$  there exists a training point  $\mathbf{x}_i$  such that

$$y_i[\mathbf{w}^\top \mathbf{x}_i + b] \leq 0$$

- There is no feasible  $(\mathbf{w}, b)$  as at least one of our constraints is violated!



## Constraints in separable setting

$$y_n[\mathbf{w}^\top \mathbf{x}_n + b] \geq 1, \quad \forall n$$

## Constraints in non-separable setting

Can we modify our constraints to account for non-separability?

Specifically, we introduce **slack variables**  $\xi_n \geq 0$ :

$$y_n[\mathbf{w}^\top \mathbf{x}_n + b] \geq 1 - \xi_n, \quad \forall n$$

- For “hard” training points, we can increase  $\xi_n$  until the above inequalities are met.
- What does it mean when  $\xi_n$  is very large? We have violated the original constraints “by a lot.”

# Soft-margin SVM Formulation

We do not want  $\xi_n$  to grow too large, and we can control their size by incorporating them into our optimization problem:

$$\begin{aligned} \min_{\mathbf{w}, b, \xi} \quad & \frac{1}{2} \|\mathbf{w}\|_2^2 + C \sum_n \xi_n \\ \text{s.t.} \quad & y_n [\mathbf{w}^\top \mathbf{x}_n + b] \geq 1 - \xi_n, \quad \forall n \\ & \xi_n \geq 0, \quad \forall n \end{aligned}$$

What is the role of  $C$ ?

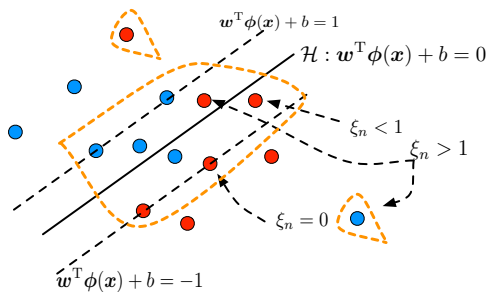
- User-defined hyperparameter
- Trades off between the two terms in our objective
- Same idea as the regularization term in ridge regression

## How to Solve this Problem?

$$\begin{aligned} \min_{\mathbf{w}, b, \xi} \quad & \frac{1}{2} \|\mathbf{w}\|_2^2 + C \sum_n \xi_n \\ \text{s.t.} \quad & y_n [\mathbf{w}^\top \mathbf{x}_n + b] \geq 1 - \xi_n, \quad \forall n \\ & \xi_n \geq 0, \quad \forall n \end{aligned}$$

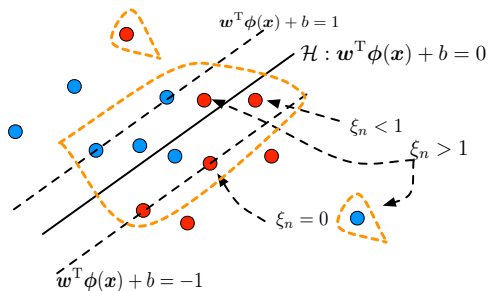
- This is a **convex quadratic program**: the objective function is quadratic in  $\mathbf{w}$  and linear in  $\xi$  and the constraints are linear (inequality) constraints in  $\mathbf{w}$ ,  $b$  and  $\xi_n$ .
- We can solve the optimization problem using general-purpose solvers, e.g., Matlab's `quadprog()` function, python's `scipy.optimize` package or CVXPY package.

## Support Vectors: Revisit



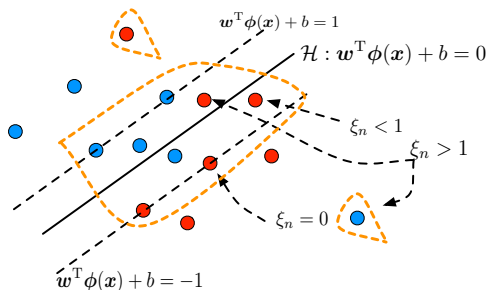
**Support vectors** are highlighted by the dotted orange lines. What does this mean mathematically?

## Support Vectors: Revisit



Recall the constraints  $y_n[\mathbf{w}^\top \mathbf{x}_n + b] \geq 1 - \xi_n$  from the soft-margin formulation. All the training points  $(\mathbf{x}_n, y_n)$  that satisfy the constraint with “=” are support vectors.

# Support Vectors: Revisit



In other words, support vectors satisfy  $y_n[\mathbf{w}^T \mathbf{x}_n + b] = 1 - \xi_n$ , which can be further divided into several categories:

- $\xi_n = 0$ :  $y_n[\mathbf{w}^T \mathbf{x}_n + b] = 1$ , the point is on the correct side with distance  $\frac{1}{\|\mathbf{w}\|}$ .
- $0 < \xi_n \leq 1$ :  $y_n[\mathbf{w}^T \mathbf{x}_n + b] \in [0, 1)$  on the correct side, but with distance less than  $\frac{1}{\|\mathbf{w}\|}$ .
- $\xi_n > 1$ :  $y_n[\mathbf{w}^T \mathbf{x}_n + b] < 0$ , on the wrong side of the boundary.

## **SVM: Hinge Loss Formulation**

---

# SVM vs. Logistic Regression

## SVM soft-margin formulation

$$\begin{aligned} \min_{\mathbf{w}, b, \xi} \quad & \frac{1}{2} \|\mathbf{w}\|_2^2 + C \sum_n \xi_n \\ \text{s.t.} \quad & y_n [\mathbf{w}^\top \mathbf{x}_n + b] \geq 1 - \xi_n, \quad \forall n \\ & \xi_n \geq 0, \quad \forall n \end{aligned}$$

## Logistic regression formulation

$$\begin{aligned} \min_{\mathbf{w}} \quad & - \sum_n \{y_n \log \sigma(\mathbf{w}^\top \mathbf{x}_n) \\ & + (1 - y_n) \log [1 - \sigma(\mathbf{w}^\top \mathbf{x}_n)]\} \\ & + \frac{\lambda}{2} \|\mathbf{w}\|_2^2 \end{aligned}$$

- Logistic regression defines a **loss for each data point** and minimizes the total loss plus a regularization term.
- This is convenient for assessing the “goodness” of the model on each data point.
- Can we write SVMs in this form as well? **The Hinge Loss formulation!**



# Derive the Hinge Loss Formulation

Here's the soft-margin formulation again:

$$\min_{\mathbf{w}, b, \xi} \frac{1}{2} \|\mathbf{w}\|_2^2 + C \sum_n \xi_n \quad \text{s.t.} \quad y_n[\mathbf{w}^\top \mathbf{x}_n + b] \geq 1 - \xi_n, \quad \xi_n \geq 0, \quad \forall n$$

Now since  $y_n[\mathbf{w}^\top \mathbf{x}_n + b] \geq 1 - \xi_n \iff \xi_n \geq 1 - y_n[\mathbf{w}^\top \mathbf{x}_n + b]$ :

$$\min_{\mathbf{w}, b, \xi} C \sum_n \xi_n + \frac{1}{2} \|\mathbf{w}\|_2^2 \quad \text{s.t.} \quad \xi_n \geq \max(0, 1 - y_n[\mathbf{w}^\top \mathbf{x}_n + b]), \quad \forall n$$

Now since the  $\xi_n$  should always be as small as possible, we obtain:

$$\min_{\mathbf{w}, b} C \sum_n \max(0, 1 - y_n[\mathbf{w}^\top \mathbf{x}_n + b]) + \frac{1}{2} \|\mathbf{w}\|_2^2$$

Divide by  $C$  and set  $\lambda = \frac{1}{C}$ , we get the **Hinge Loss formulation**:

$$\min_{\mathbf{w}, b} \sum_n \underbrace{\max(0, 1 - y_n[\mathbf{w}^\top \mathbf{x}_n + b])}_{\text{Hinge Loss for } x_n, y_n} + \frac{\lambda}{2} \|\mathbf{w}\|_2^2$$


# Cross-Entropy Loss vs. Hinge Loss

Given training data  $(x_n, y_n)$ , the cross entropy loss was

$$-\{y_n \log \sigma(\mathbf{w}^\top \mathbf{x}_n) + (1 - y_n) \log[1 - \sigma(\mathbf{w}^\top \mathbf{x}_n)]\}$$


$$-\log \left( \frac{e^{-\mathbf{w}^\top \mathbf{x}_n}}{1 + e^{-\mathbf{w}^\top \mathbf{x}_n}} \right)$$

If  $y = 0$



$$-\log \left( \frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}_n}} \right)$$

If  $y = 1$

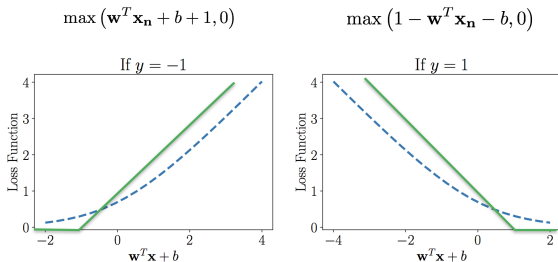


What does the Hinge Loss Function look like?

# Cross-Entropy Loss vs. Hinge Loss

Given training data  $(\mathbf{x}_n, y_n)$ , the Hinge loss is

$$\max(0, 1 - y_n[\mathbf{w}^\top \mathbf{x}_n + b])$$



- Loss grows linearly as we move away from the boundary.
- No penalty if a point is more than 1 unit from the boundary.
- Makes the search for the boundary easier (as we will see later).

# Hinge Loss SVM Formulation

Minimizing the total hinge loss on all the training data

$$\min_{\mathbf{w}, b} \sum_n \underbrace{\max(0, 1 - y_n[\mathbf{w}^\top \mathbf{x}_n + b])}_{\text{hinge loss for sample } n} + \underbrace{\frac{\lambda}{2} \|\mathbf{w}\|_2^2}_{\text{regularizer}}$$

Analogous to regularized least squares or logistic regression, as we balance between two terms (the loss and the regularizer).

- Can solve using gradient descent to get the optimal  $\mathbf{w}$  and  $b$
- Gradient of the first term will be either 0,  $\mathbf{x}_n$  or  $-\mathbf{x}_n$  depending on  $y_n$  and  $\mathbf{w}^\top \mathbf{x}_n + b$ .
- Much easier to compute than in logistic regression, where we need to compute the sigmoid function  $\sigma(\mathbf{w}^\top \mathbf{x}_n + b)$  in each iteration.

# Three SVM Formulations

**Hard-margin (for separable data)**

$$\min_{\mathbf{w}, b, \xi} \frac{1}{2} \|\mathbf{w}\|_2^2 \text{ s.t. } y_n[\mathbf{w}^\top \mathbf{x}_n + b] \geq 1, \xi_n \geq 0, \forall n$$

**Soft-margin (add slack variables)**

$$\min_{\mathbf{w}, b, \xi} \frac{1}{2} \|\mathbf{w}\|_2^2 + C \sum_n \xi_n \text{ s.t. } y_n[\mathbf{w}^\top \mathbf{x}_n + b] \geq 1 - \xi_n, \xi_n \geq 0, \forall n$$

**Hinge loss (define a loss function for each data point)**

$$\min_{\mathbf{w}, b} \sum_n \max(0, 1 - y_n[\mathbf{w}^\top \mathbf{x}_n + b]) + \frac{\lambda}{2} \|\mathbf{w}\|_2^2$$

You should know:

- Max-margin formulation for separable and non-separable SVMs.
- Definition and importance of support vectors.
- Hinge loss formulation of SVMs.
- Equivalence of the max-margin and hinge loss formulations.